D123 | Update beliefs, not old predictions

PRIVATE FICTIONAL firm-growth forecast. No actual firm data, investment, trade, research purchase, founder contact, public prediction, disclosure or spending.

Event before evidence	At least30% same-definition audited ARR growth in next12months from fixed baseline. Fix currency, acquisition/price effects, source, cutoff and resolution rule. This workbook has no real firm-window data or validated reference-class selection.
Outside-in update	20 successes/100 completed comparable windows is an assumed20% prior. Signal conditional rates.6 success/.3 failure give joint.12/.24, posterior1/3. Likelihood ratio2 is not posterior60%.
Duplicate is not new evidence	Second newsletter repeats exactly the same audit, no new information: posterior stays1/3. Naively multiplying LR2 twice gives50%. This wrong-counterfactual is labeled, not a supported update. Independent repeat signals require conditional independence evidence.
Revised prior	Separate scenario prior10% with same signal gives2/11=18.181818%. Do not secretly mix a changed prior with a copied posterior. Boundary priors0/1 and impossible signal are explicit; odds can be INFINITE or UNDEFINED.
Frozen scoring	Four different predeclared resolved events: p.8/.7/.3/.2 versus y1/0/1/0. Binary average squared error.265; locked benchmark.5 gives.25, difference.015. Two-class-coordinate sum is twice the score. This tiny ledger is not population calibration proof.
Version and authority	Synthetic lock copies and day offsets only check internal consistency, not tamper-proof histories. Blank unresolved outcomes hold aggregate rather than become failures. Static prose is Case-v1. Actual learner/error/reviewer/transfer fields blank; no overall10 claim.

Inputs | Assumed class and signal

Synthetic reference-class counts and conditional signal rates. No real company, security, diagnostic decision or validated predictive model. All yellow values editable.

Reference-class completed windows	100.00
Reference-class successes	20.00
Signal probability conditional on success	0.60
Signal probability conditional on failure	0.30
Separately revised relevant prior	0.10

Reference/model valid? (1=yes)	1.000000
Revised prior valid? (1=yes)	1.000000

Reference-class selection, completed-window eligibility and conditional rates require separate evidence in a real forecast. Two reports copying one audit do not establish independent likelihoods. These checks are arithmetic bounds, not empirical or account clearance.

Bayes | Joint denominator before odds

Same event and observed signal, conditional likelihoods fixed. Impossible signal posterior UNDEFINED. The wrong duplicate comparison is a teaching counterfactual, never a recommended update.

Prior success probability	0.200000
Prior success odds	0.250000
Likelihood ratio of signal	2.000000
Joint success and signal	0.120000
Joint failure and signal	0.240000
Signal marginal probability	0.360000
Posterior success probability	0.333333
Posterior success odds	0.500000
Duplicate audit: unchanged probability	0.333333
WRONG independent-repeat assumption	0.500000
Separate revised-prior posterior	0.181818

Exact duplicate contains literally no additional information. New evidence may matter, but independence cannot be inferred from different URLs, analysts or publishers. A prior0 with a signal impossible under that prior is a model conflict, not automatic success.

Ledger | Synthetic frozen forecasts

Four DIFFERENT resolved binary events. Day0 forecast before day365 resolution; offsets are fictional, not real timestamps. A lock-copy equality is a consistency check, not tamper-proof provenance.

Event	p now	Outcome y	Locked p	Issue day	Resolve day	Row valid	Squared loss	Bench .5 loss
F1	0.80	1.00	0.80	0.00	365.00	1.000000	0.040000	0.250000
F2	0.70	0.00	0.70	0.00	365.00	1.000000	0.490000	0.250000
F3	0.30	1.00	0.30	0.00	365.00	1.000000	0.490000	0.250000
F4	0.20	0.00	0.20	0.00	365.00	1.000000	0.040000	0.250000

Blank or ambiguous outcome remains unresolved under its predeclared rule. This fixed complete-ledger evaluation holds aggregate scoring if any row is unresolved/invalid; it never drops a row to improve the score. Editing both current and lock-copy can evade equality: real immutable version evidence still needed.

Score | State the Brier convention

Equal weight, same four events and horizon. Binary score = mean (p-y)^2 in[0,1]. Two class-coordinate sum ranges[0,2]. Do not compare differently scaled unlabeled scores.

All four rows ready? (1=yes)	1.000000
Binary mean squared-error Brier	0.265000
Precommitted constant .5 benchmark	0.250000
Difference, positive means worse than benchmark	0.015000
Two-coordinate sum mean, different scale	0.530000
Two-coordinate benchmark	0.500000
This ledger comparison only	HIGHER LOSS

Brier mixes calibration, resolution and uncertainty. Lower overall loss is not calibration-only proof. Four outcomes cannot establish population calibration, a universal .5 strategy or that every confident prediction is bad. Do not rewrite forecasts after resolution and call it improvement.

Bin | Small descriptive agreement

Independent ten-event thought experiment, not the four-row ledger. Count8/10 versus forecast.8 is empirical agreement in one tiny bin. No population confidence or resolution claim.

Number of same-bin forecasts	10.00
Number of observed successes	8.00
Common forecast probability	0.80

Bin inputs valid? (1=yes)	1.000000

Observed bin event fraction	0.800000
Observed fraction minus forecast	0.000000

8/10 is not eight guaranteed future successes. These counts do not establish population calibration or between-bin discrimination. A single failure at high probability is not by itself evidence the probability was incorrect. No automatic extremizing rule is adopted.

Event rules | Fix resolution in advance

Proposed fictional day0/day30 workflow, not a real investment committee or permission grant. No real owner accepted these roles.

Event identity	One eligible firm/window per event, fixed baseline and12-month horizon. Specify audited recurring revenue definition, currency/FX, acquisition effects and price changes before forecasting. The four score rows represent different events, not four revisions of one event.
Evidence role | day0	Record reference-class inclusion/exclusion, eligible completed denominators, prior version, signal source/conditional rates and dependency map. Link duplicate newsletters to their original audit. Unknown reference-class validity remains unknown.
Forecast role | day0	Keep actual issue/update timestamps, probability and reason before resolution. Any genuine prior revision gets its own version. Lock original predictions and precommit benchmark/convention; synthetic copies here cannot prove real timestamp integrity.
Review role | day30	Review objections, dependency and version consistency. Outcome remains unresolved until the12-month event resolves under stated source/cutoff rule. Day30 is not early final resolution. No silently changed definitions or selective deletion.
Hold gate	Unknown event/source/authority blocks real use. Missing or ambiguous outcomes stay unresolved; this fixed ledger holds the aggregate. A real policy for eligible resolved-subset scoring would need predeclaration and disclosure, not retrospective cherry-picking.
Delayed transfer	Closed notes: revise prior to.1 and rederive2/11; explain why repeated audit is not another LR. Freeze a new fictional forecast set before outcomes. Actual performance/error/objection/transfer records stay blank until observed.

Evidence | Written sources, bounded claims

Selected written source text rechecked13 November2026. Not full talk, audiovisual continuity, source diagrams, empirical venture base rates, real learner testing or independent validation.

Penn State STAT509	https://online.stat.psu.edu/stat509/lesson/17/17.3 - prior/prevalence, conditional predictive values and LR definitions. Clinical examples are not startup base rates; no medical application is being performed.
scikit-learn calibration	https://scikit-learn.org/stable/modules/calibration.html - bin fractions versus predictions; Brier combines calibration/resolution/uncertainty; fitting calibration on held-out data matters. No classifier trained or calibrated here.
Long Now written summary	https://longnow.org/talks/02015-tetlock/ - outside-in numerical forecasts and keeping score. Written talk summary is not the full recording. Historical tournament30% claim not independently verified or adopted as an expected gain.
Corrections and limits	One matching small bin is not population calibration or resolution proof. No automatic extremizing from group agreement. Separate binary single-coordinate from two-coordinate score conventions; compare identical horizons/eligibility/weights.
Strongest objection	Forecasts may be internally valid but based on a poor reference class or dependent evidence. Arithmetic cannot settle that. Preserve assumption/source/dependency versions and actual disagreement rather than invent external validation or a profitable investment.

Actual work | No observed forecast or learner record

Record only actual work, errors, revisions, reviewer words and changed-context transfer. All fields blank; synthetic forecasts are not public predictions or realized investment results.

Actual date/version	Prior/evidence source	Original response/error	Dependency correction	Frozen ledger proof	Objection/changed view	Delayed transfer/result

Actual work remains blank. Record only what occurred: dates, reasoning, error/source/correction, strongest objection and a new delayed transfer. No invented learner results, sponsor acceptance, actual budget release or signed rights.

Failure log | Authored execution only

Observed native tests are added only after execution. Not actual learner mistakes, independent human audit, Excel/Sheets parity, real authority or overall 10/10.

Controlled test	Observed authored outcome
Baseline Bayes/odds/duplicate	PASS:all baseline Bayes/odds/duplicate formulas independently checked.
Boundary/changed priors and signals	PASS:prior0/1/revised priors/zero and perfect signals.
Changed frozen scores/benchmark	PASS:24 changed sets and restoration; scoring extremes/tie/inverted outcomes.
Missing/invalid ledger holds	PASS:17 invalid cases; missing outcome holds aggregate, not failure.
Version/time conflicts	PASS:lock mismatch and issue/resolve conflicts held; no immutable proof.
Bin boundaries/invalid counts	PASS:bin0/all successes and forecast endpoints, invalid counts.
Broken result reference hold	PASS:7 broken critical references held affected outputs.
Restored baseline/actual blank	PASS:baseline restored; actual learner/reviewer/transfer fields blank.

Only the original synthetic course exercise at https://top1os.dev/ supplies the case figures. No new empirical reading, real participants, company records, Indian legal clearance, actual commitments or source completeness claimed.