D 150 | Coverage is not investment alpha

PRIVATE FICTIONAL offline sourcing lab. Scores / counts / lakh INR labeled. No actual company records, paid feed, scraper, API, private account connection, outreach or spend.

Original score toy	Weights 20/20/25/20/15 and normalized signals .8/.6/.4/.9/.2 -> weighted numerator 59 at weight sum 100, full score 59/100. It is an unvalidated ranking score, not 59% success probability or investment return.
Missingness is not zero	Separate availability flags 1/1/0/1/1 exclude third signal from partial score. Available weighted numerator 49, covered weight 75 -> normalized score 65.333333/100, coverage 75%. Higher normalized score is not proved better quality. Complete and partial results describe different observability.
Complete-label screening toy	TP 30 / FP 70 / FN 10 / TN 390 =500 eligible companies. Flagged 100 contains 30 relevant; unflagged 400 contains 10 relevant. Precision 30%, recall 75%, base rate 8%; classification labels assumed complete, correct, comparable and same window. No real sourcing coverage or alpha inferred.
Original cost toy	Build one-off 8 + annual licenses 6/run 6/security 2 ->22 year one, 14 recurring. Buy annual 10 + one-off integration 2 + annual operations 4 ->16 year one, 14 recurring. Analyst annual 18 is stipulated, not a quote or equivalent role. No automatic salary saving.
Separate cost horizon	Three undiscounted full years at unchanged recurring rates: build 8+3*14=50, buy 2+3*14=44. Difference 6 is fixed one-off gap here, not a universal recommendation. No inflation, taxes, financing, discounting, benefit, implementation delay, contract cancellation or actual purchase.
Design before activation	Resolve identity/provenance, permission/purpose, as-of feature and collection times, group separation, label delay and blind spots. Two splitters alone do not jointly enforce time+family separation. Actual evidence/authority/accepted owner stays blank. No learner, account, legal/field or overall 10 validation.

Signals | Score and observation are separate

Shaded inputs: weights B, signal C in[0, 1], available D=1/0. E complete gate, F partial gate. D 0 excludes C only from partial case; it does not assert observed zero. Labels synthetic.

Synthetic signal	Weight	Normalized value	Available 1 / 0	Complete gate	Partial gate
Signal 1	20.00	0.80	1.00	1.000000	1.000000
Signal 2	20.00	0.60	1.00	1.000000	1.000000
Signal 3	25.00	0.40	0.00	1.000000	1.000000
Signal 4	20.00	0.90	1.00	1.000000	1.000000
Signal 5	15.00	0.20	1.00	1.000000	1.000000

All five weights nonnegative; sum may be edited, outputs explicitly normalize by actual weight sum. No negative weights or signal outside 0-1. Partial row gate permits C blank/text when D 0 because excluded, not silently made an observed zero. With D 1 it must be numeric 0-1. No probability calibration, selection-bias fix or data permission follows.

Inputs | Independent labels and costs

All fictional. Counts nonnegative integers; costs lakh INR nonnegative; horizon positive whole full years. Complete labels and unchanged annual rates stipulated, not observed.

True positives: flagged and relevant	30.00
False positives: flagged, not relevant	70.00
False negatives: unflagged, but relevant	10.00
True negatives: unflagged, not relevant	390.00

Build one-off implementation	8.00
Build annual licenses	6.00
Build annual run / operations	6.00
Build annual security	2.00
Buy annual subscription	10.00
Buy one-off integration	2.00
Buy annual operations	4.00
Stipulated analyst annual cost; not equivalent role	18.00
Full-year undiscounted comparison horizon	3.00

Complete-label count gate	1.000000
Cost numeric gate	1.000000

Zero-count denominators undefined, not observed 0 performance. Analyst cost enters only a labeled cash-cost difference, not salary savings. No actual prices, comparable coverage, staff capacity, security compliance, benefit, ROI or spending authority. Static baseline prose must be revised before exporting edited cases.

Scores | Normalize without inventing data

Conditional fictional arithmetic. Undefined denominators stay #N/A. Independent families do not validate source coverage, permissions or business outcomes.

All weight sum	100.000000
Full weighted numerator	59.000000
Full normalized score /100	59.000000

Partial total weight	100.000000
Covered weight	75.000000
Available weighted numerator	49.000000
Partial normalized score /100	65.333333
Covered fraction of all weight	0.750000
Missing fraction of all weight	0.250000
Partial less full score; not quality improvement	6.333333

Baseline full 100/59/59; partial 100/75/49/65.333333/.75/.25; score difference 6.333333 not evidence of improvement. If all flags 0, total weight known, coverage 0 but normalized partial scoreundefined. All weights 0 makes scores / coverage undefined. Invalid ignored C cannot invalidate partial when D 0; it still invalidates full. No ranking comparability or calibrated success probability asserted.

Screening | Count the missed positives

Conditional fictional arithmetic. Undefined denominators stay #N/A. Independent families do not validate source coverage, permissions or business outcomes.

Eligible total	500.000000
Flagged / review workload	100.000000
Unflagged total	400.000000
All relevant	40.000000
Precision: TP / flagged	0.300000
Recall: TP / relevant	0.750000
Relevance base rate	0.080000
Flagged fraction	0.200000
Miss fraction among relevant	0.250000
False-positive fraction among flagged	0.700000
Population partition identity; must be zero	0.000000

Baseline 500/100/400/40/.3/.75/.08/.2/.25/.7/0. All 500 labels assumed complete / correct, not inferred from only invested/reviewed deals. If unflagged labels unknown, FN/TN must not be fabricated as 0; recall/base rate unestablished. Confusion counts do not verify unseen-universe coverage, independent labels, drift, capacity, returns or investment alpha.

Costs | Price is not role equivalence

Conditional fictional arithmetic. Undefined denominators stay #N/A. Independent families do not validate source coverage, permissions or business outcomes.

Build annual recurring	14.000000
Buy annual recurring	14.000000
Build year one total	22.000000
Buy year one total	16.000000
Build minus buy, year one cost	6.000000
Build horizon cost, undiscounted	50.000000
Buy horizon cost, undiscounted	44.000000
Build minus buy, horizon cost	6.000000
Build year one minus analyst annual cost	4.000000
Buy year one minus analyst annual cost	-2.000000
Recurring build minus buy	0.000000

Baseline 14/14/22/16/6/50/44/6/4/-2/0 at three full years. Analyst differences are cash-cost comparisons, not job savings or benefits. Recurring equal does not mean quality/returns equal. Annual sums linear, not NPV; one-off counted once. No supplier quotes, license terms, implementation delay, taxes, inflation, discounting, ROI or approved purchase.

Identity and rights | Architecture without collection

No scraping script, extension, API, feed export, private-network sync, real founder data, outreach or purchase. These are proposed design checks only.

Stable identity	Use a stable synthetic company ID with versioned legal/domain/person aliases, provenance and match confidence. Preserve conflicts; route uncertain matches for review rather than merge solely on name/city. Resolve entity-family aliases before splitting to avoid duplicate leakage.
Source authority	Record source, license/access/purpose, allowed fields, audience, retention, correction and refresh evidence before real collection. Public visibility or vendor marketing is not bulk collection, redistribution or private data permission. Company-level UPI/GST revenue not inferred from aggregate dashboards.
LinkedIn boundary	Official returned help prohibits third-party scraping/automation, fake accounts, security/access bypass and unconsented copying/distribution. Do not install a crawler, extension, connect private accounts or contact founders from this classroom design. No legal opinion about another source license.
Vendor claims	Ensemble Unity January 3, 2025 written page uses strong coverage/superiority language, not independent proof. Harmonic returned homepage advertises features, not verified India coverage, refresh SLA, prices or alpha. No demo request, purchase, signup, feed/API or inbox/calendar/network connection.
Prediction is not guarantee	Retterath February 26, 2026 written 2030 thought experiment is advocacy/prediction, not measured inevitable model improvement or salary savings. Biased historical passes and selective outcomes can reinforce errors. Label gaps and drift require monitoring evidence; no actual model training here.
Stop condition	Unknown ownership, permission, source reliability, entity match, time availability or label definition pauses affected real collection/use. Actual evidence, authority and accepted owner remain blank. No real person-company dossier or participant data required.

Validation design | Time and group both matter

Proposed offline design, not a trained/running system or validated investment outcome. No universal best signal or three-to-six-month window established.

As-of availability	Store event time, first observed/collected time, source/version and correction history. An event predating a score cutoff is still unavailable if first collected afterwards. Do not backfill later financing/outcome/news into earlier features as if known then.
Freeze before test	Define target, horizon, eligibility, decision cutoff, label maturity, threshold, capacity and metrics before evaluation. Split before learning transforms/feature selection/imputation/scaling; fit on training data only, apply same learned transforms to test. Do not tune on reported test outcomes.
Different split guarantees	TimeSeriesSplit protects chronological order, not duplicate/company-family separation. GroupKFold protects nonoverlapping groups, not chronological availability. Combining goals needs an explicit design satisfying both; calling either alone does not establish both.
Irregular records	TimeSeriesSplit documentation requires equally spaced samples for comparable fold-duration metrics; gap is a number of samples, not automatically calendar days or label delay. Irregular company events need justified time windows and maturity gaps. No splitter applied or real sample count claimed.
Blind spots and workload	Record eligible-universe/source coverage, missingness by subgroup, unflagged blind-spot samples, label verification, reviewer capacity and overrides. Unknown unflagged outcomes are not 0 false negatives. Scores/precision alone do not establish causal value or returns; invested-deals-surfaced-first alone not ROI.
Closed-notes transfer	Set third signal available 1: score 59 / coverage 100%; then restore. Set TP 0 with FP 70 / FN 10 / TN 390: precision 0 / recall 0 with positive denominators. Set all flags 0: coverage 0, partial score undefined. Raise buy annual subscription 2 lakh: year one 18, three-year 50, not 16/44. Actual learner answers blank.

Evidence | Documentation versus vendor claims

Nine URLs/five organizations: software documentation, platform policy, vendor marketing and practitioner prediction. Not independent investment-alpha studies or source completeness.

Metrics definitions	https://scikit-learn.org/stable/modules/generated/sklearn.metrics.precision_score.html and https://scikit-learn.org/stable/modules/generated/sklearn.metrics.recall_score.html | Selected definition and zero_division notes read, not all examples/API. Library may return 0 with warning/configuration; lab intentionally uses undefined #N/A when denominator 0.
Leakage and split scope	https://scikit-learn.org/stable/common_pitfalls.html | Selected 12.2/12.2.1 read, not whole page. https://scikit-learn.org/stable/modules/generated/sklearn.model_selection.TimeSeriesSplit.html and https://scikit-learn.org/stable/modules/generated/sklearn.model_selection.GroupKFold.html | Selected opening/parameters read, not all examples/plots or a custom combined splitter.
Platform policy	https://www.linkedin.com/help/linkedin/answer/a1341387 | Complete returned written help read: prohibited scraping/automation/fake accounts/bypass/unconsented redistribution. No actual account restriction, login, connection, automation or permission review performed.
Vendor descriptions	https://www.ensemble.vc/research/unity-data-driven-sourcing | Returned written page read, January 3, 2025, not diagram / dataset validation. https://www.harmonic.ai | Returned homepage written overview, not images/linked resources / customer proof / India-specific coverage/prices/terms. Advertising not observed performance or authority.
Practitioner thesis reused	https://www.newsletter.datadrivenvc.io/p/do-you-still-need-analysts-the-investment-firm-of-2030-fewer-people-better-decisions | Earlier complete returned written thought experiment reused, visible February 26, 2026, not images / underlying datasets/advertised guide. Prediction not validated productivity, guaranteed learning, investment alpha or job-saving proof.
Outcome limits	Original figures synthetic. No real company/source/license/label/alpha study, collection/API/feed/account/outreach/spend, independent learner, native Excel/Sheets parity, full source/AV or overall 10. Actual evidence/authority/accepted owner plan blank. No pipeline service installed, trained, activated or tested.

Actual plan | No real data or activation

Blank actual source/right/coverage/as-of/group/label/authority record. Proposed reviewer roles are not actual accepted owners. No startup list or private contact data supplied.

Source / right	Entity / match	As-of / cutoff	Group / labels	Coverage / gap	Owner / authority	Stop / review

No actual data purchase, source access, model training, collection, account connection, extension, API, private-network sync, outreach, learner test or spend. Blank is unknown, not license granted, zero missingness, zero cost or accepted ownership. Current vendor coverage/prices/alpha unverified.

Failure log | Authored arithmetic only

No observed real sourcing system, company outcomes, vendor alpha, learner attempt, account flow or overall 10.

1	PASS: 44 baseline formulas independently recomputed.
2	PASS: 22 changed input sets plus restored baseline.
3	PASS: 24 invalid inputs/1056 family-isolation checks.
4	PASS: 6 broken references/6 checks.
5	PASS: 8 boundaries;missing/zero weights/counts/denominators.
6	PASS: score/full-partial/labels/cost families isolated.
7	PASS: ignored missing signal not observed zero.
8	PASS: actual evidence/authority/accepted-owner rows blank.