D 155 | Forecast history is not hindsight

PRIVATE FICTIONAL offline journal. No public forecasts, actual outcomes, reference calls, investment recommendation, account or monitoring.

Fixed checkpoint	Four resolved toy questions have original p .8/.6/.3/.1 and y 1/0/1/0. Binary Brier mean .225 versus fixed .5 baseline .25; relative skill .10. All probabilities use 0-1, lower loss better; this is not accuracy or a calibration proof.
Appended alternative	Updates .9/.4/.6/.2 on those same four outcomes give .0925 and relative skill .63. Difference .1325 is descriptive, not evidence of human improvement. Updates need their own prespecified checkpoint and must precede resolution; this numeric toy does not verify timestamps.
Eligibility before averaging	Eight toy rows: four resolved, two unresolved, two void. Unresolved and void outcomes blank, excluded from all three paired scores. Status codes 1 resolved, 0 unresolved, 2 void. Never encode unknown as no or drop bad resolved questions after seeing outcomes.
Preserve history	Original probability and rationale stay. Append timestamped update/evidence, do not overwrite. Workbook editing is not a tamper-proof audit log. No automatic time weighting, resolver, event retrieval, rescheduling, actual calibration fit or future monitor.
References separate	Eight fictional role slots and four blank interview-note blocks. Common core plus context follow-up, permission/audience and firsthand examples. Pauses or nonanswers do not establish negative findings. No actual contact, quotes, consent or accepted owner fabricated.
Transfer without notes	Set resolved Q 2 original p to .9: original loss .3375/skill -.35, alternative unchanged .0925. Set Q 4 status unresolved and clear outcome: eligible 3, original .296666667, alternative .11, baseline .25. Static prose is baseline only; revise before changed export.

Toy inputs | Eight synthetic rows

No real event or timestamp. Pairs use one common eligible set. p, alternative and fixed reference must be numeric 0-1.

Toy ID	Original p	Alternative p	Reference p	Outcome y	Status 0/1/2	Valid row?
Q1	0.80	0.90	0.50	1.00	1.00	1.000000
Q2	0.60	0.40	0.50	0.00	1.00	1.000000
Q3	0.30	0.60	0.50	1.00	1.00	1.000000
Q4	0.10	0.20	0.50	0.00	1.00	1.000000
Q5	0.70	0.80	0.50		0.00	1.000000
Q6	0.20	0.30	0.50		0.00	1.000000
Q7	0.90	0.80	0.50		2.00	1.000000
Q8	0.40	0.50	0.50		2.00	1.000000

Outcome required only for resolved rows, binary 0/1. Unresolved/void requires truly empty outcome; zero is an observed no, not missing. Invalid row gives undefined aggregate instead of silent dropping. Resolved does not authenticate a source or prior timestamp.

Row losses | Same eligible questions

Blank loss means valid but unresolved/void, not zero loss. Invalid row yields an error, not favorable exclusion.

Toy ID	Eligible 1/0	Original loss	Alternative loss	Reference loss
Q1	1.000000	0.040000	0.010000	0.250000
Q2	1.000000	0.360000	0.160000	0.250000
Q3	1.000000	0.490000	0.160000	0.250000
Q4	1.000000	0.010000	0.040000	0.250000
Q5	0.000000			
Q6	0.000000			
Q7	0.000000			
Q8	0.000000			

Baseline original .04/.36/.49/.01; alternative .01/.16/.16/.04; reference .25 each. Four excluded rows blank, not y=0. Pairing requires all forecast columns valid so denominators match. This strict common-set policy is a declared choice, not every platform rule.

Summary | Loss, skill and descriptive gaps

All eight rows must pass numeric/status gate; four eligible in baseline. No significance test, calibration estimator or investment return.

All rows valid (1=yes)	1.000000
Eligible resolved count	4.000000
Original binary Brier loss	0.225000
Alternative binary Brier loss	0.092500
Reference loss on same rows	0.250000
Original skill 1 - loss/reference	0.100000
Alternative skill 1 - loss/reference	0.630000
Original minus alternative loss	0.132500
Original mean p, same eligible rows	0.450000
Observed yes fraction, eligible only	0.500000
Original mean p minus observed yes fraction	-0.050000
Unresolved count	2.000000
Void count	2.000000

Baseline n 4, .225/.0925/.25, skill.10/.63, loss gap.1325, mean p.45/yes.50/gap-.05. Aggregate gap is not full calibration. Skill may be negative; reference loss zero makes relative skill undefined. No resolved rows leaves loss undefined. Brier is proper; do not assume every derived skill ratio is strictly proper.

History policy | Define before seeing outcomes

Offline practice only. The input sheet is an editable calculator, not a timestamped journal or actual resolution record.

Definition and resolver	For each of 20 blank slots specify observable yes/no event, deadline/timezone, source/resolver, data cutoff, amendment/cancellation/missing-evidence rules and void criteria. Avoid ambiguous fundraising rumors. Future claims are not current facts.
Fixed checkpoint	Choose one timestamp or lead-time rule before outcomes. Only probabilities genuinely available before that checkpoint qualify. Preserve original evidence and rationale. Do not select each question best historical forecast after resolution.
Updates	Append update timestamp, p, new evidence and reason. Late information can improve descriptive score without proving comparable skill. If time-weighted scoring is wanted, define time intervals/weights first; this workbook does not implement it.
Common-set comparison	Here, all original/alternative/reference probabilities must be present and valid, and y binary for resolved rows. One invalid row blocks aggregate. Valid unresolved/void is excluded symmetrically. Real exclusions and late forecasts require prespecified rules, not hindsight.
Calibration and fitting	Lower Brier can reflect reliability, resolution or uncertainty. Four observations/20 blank slots do not estimate general calibration. Do not fit an adjustment and claim out-of-sample improvement on the same outcomes. Holdout/temporal split and sufficient comparable data matter; no universal sample threshold.
Scope and hypothesis	9/10 confidence is not automatically p=.9. Compare relevant event populations and baselines. A premortem creates hypotheses, not measured failure probabilities or bias immunity. Cheapest information is not always decision-relevant. No calendar or monitoring created.

Journal 20 | Blank prospective design slots

All fields blank. No fabricated forecast, outcome, timestamp, resolver, consent or automatically scheduled review.

Slot	Event / yes rule	Deadline / zone	Resolver / void rule	Original p / time	Rationale / source	Append update / time
1						
2						
3						
4						
5						
6						
7						
8						
9						
10						
11						
12						
13						
14						
15						
16						
17						
18						
19						
20						

For real work, add outcome/source/resolution time as an appended record. The 20 empty slots are not 20 completed or scored forecasts. No actual horizon, monitor, reference call, account or public post is authorized by filling this design sheet.

Reference protocol | Draft questions, not contacts

Eight fictional role slots: two former bosses, two former reports, two cofounders, two customers. No actual people identified or approached.

Role slots	Record relationship, time span, firsthand access, selection bias, permission, audience and limits. Eight slots is this exercise design, not required sample size or permission for backchannel contact. CRV source is founders checking investors, not equivalent empirical founder-check evidence.
Shared core	What concrete decision did you observe? What happened under pressure? What did the person commit to and do? What counterexample changes your view? Ask only within known relationship, access and authorized scope.
Follow-up and caution	Adapt context and bounded concern questions. They cannot guarantee candor. A pause, hedge, politeness or refusal of financial opinion is not a negative verdict. Separate firsthand example, hearsay, inference and unknown.
Draft note 1	Relationship/access: blank. Permission/audience: blank. Observed example/exact words/source date: blank. Counterevidence/unknowns: blank. No interview occurred.
Draft note 2	Relationship/access: blank. Permission/audience: blank. Observed example/exact words/source date: blank. Counterevidence/unknowns: blank. No interview occurred.
Draft note 3	Relationship/access: blank. Permission/audience: blank. Observed example/exact words/source date: blank. Counterevidence/unknowns: blank. No interview occurred.
Draft note 4	Relationship/access: blank. Permission/audience: blank. Observed example/exact words/source date: blank. Counterevidence/unknowns: blank. No interview occurred or disclosure authorized.

Actual evidence | Unknown is not no

Blank evidence and permission stay blank. No public forecast, real reference, account, investment commitment, monitoring or spending.

Question / scope	Source / date	Observed response	Error / uncertainty	Permission / audience	Accepted owner	Next review / authority

Actual learner answers, reference interviews, resolution evidence and consent not present. Do not convert a day score into question-level diagnosis. Authored calculator tests are not independent learner or live account validation. Full source AV/native Excel-Sheets parity/actual overall 10 remain open.

Evidence | Written sources and explicit limits

Six URLs/five organizations/four types: official methods, practitioner, abstract, community summary. Not six independent studies.

Official loss	https://scikit-learn.org/stable/modules/generated/sklearn.metrics.brier_score_loss.html | Complete returned written API docs read, not linked code/gallery/papers. Binary convention 0-1 differs from unhalved/multiclass 0-2. Native package version shown by source not an installed parity test.
Official calibration	https://scikit-learn.org/stable/modules/calibration.html | Complete returned written docs read, not figure pixels/underlying papers. Brier mixes calibration, resolution, uncertainty; disjoint fitting data matters. Example 1000 sample isotonic advice not universal human forecasting threshold.
Practitioner references	https://www.crv.com/content/reference-check-questions-for-founders | Complete returned April 21, 2026 written read. Founders checking investors. Hedged-answer negative heuristic and five to seven calls are author views, not validated universal signals or permission. No outreach/backchannel routine followed.
Abstract only	https://papers.ssrn.com/sol3/papers.cfm?abstract_id=4818073 | Complete returned abstract/metadata only; written Feb 10, 2024, posted May 7/revised May 9, 2024. Not 63 page paper/underlyingstudies. Fifteen-bias abstract claim not independently validated or effect size/expert hit-rate proof.
Tournament notebook	https://www.metaculus.com/notebooks/43356/calibration-adjustment-analysis/ | Complete returned written/table read, not figures/code/data/replication. 2025 Q 1-Q 2 bot dataset; 500 forecast filter heuristic, temporal 70/30 split. Reported improvements not guaranteed for human 20 slots; multiclass per-option normalization not our binary rule. Not current platform points.
Community baseline	https://forecasting.wiki/wiki/Brier_score | Complete returned written summary read; it says imported live from Wikipedia, not independent primary validation. Formula 1-loss/reference used with fixed reference. Its lower loss implies better calibration wording conflicts with official decomposition caution; not adopted. No underlying citations, chart pixels, actual calibration fit/full source AV/independent learner/account/overall 10.

Failure log | Authored checks only

Independent arithmetic and controlled changes are not real forecast skill, reference outcomes, learner results or account tests.

1	PASS: 53 baseline formulas independently recomputed.
2	PASS: 20 changed sets;original and alternatives isolated.
3	PASS: 20 invalid sets/1060 checks.
4	PASS: 8 boundaries;no resolved and zero reference loss handled.
5	PASS: 6 broken result references propagate errors.
6	PASS: unresolved/void loss blank,not zero or no.
7	PASS: blank 20-slot journal and actual evidence retained.
8	PASS: restored baseline,not native cross-viewer/learner/actual skill.